

ESCUELA TÉCNICA SUPERIOR DE  
INGENIERÍA INFORMÁTICA

MÁSTER UNIVERSITARIO EN INGENIERÍA  
DEL SOFTWARE E INTELIGENCIA ARTIFICIAL

Diseño de un método comparativo para copilotos de  
código con IA

Design of an Evaluation Method for AI  
Programming Assistants

Realizado por  
Donat Shergalis

Tutorizado por  
Gabriel Jesús Luque Polo

Departamento  
ESCUELA TÉCNICA SUPERIOR DE  
INGENIERÍA INFORMÁTICA,  
UNIVERSIDAD DE MÁLAGA

MÁLAGA, Septiembre de 2025







ESCUELA TÉCNICA SUPERIOR DE INGENIERÍA INFORMÁTICA  
GRADO EN MÁSTER UNIVERSITARIO EN INGENIERÍA DEL SOFTWARE E  
INTELIGENCIA ARTIFICIAL

Diseño de un método comparativo para copilotos de código con IA  
Design of an Evaluation Method for AI Programming Assistants

Realizado por  
**Donat Shergalis**

Tutorizado por  
**Gabriel Jesús Luque Polo**

Departamento  
**ESCUELA TÉCNICA SUPERIOR DE INGENIERÍA INFORMÁTICA**

UNIVERSIDAD DE MÁLAGA  
MÁLAGA, Septiembre de 2025

Fecha defensa:  
El Secretario del Tribunal



## Agradecimientos

Página de agradecimientos.



**Resumen:**

Resumen ES.

**Palabras claves:** clave1, clave2, ..., clave5

---

**Abstract:**

Benchmarks are widely used for measuring and comparing performance of LLMs in different areas, including logical puzzles, factual accuracy, and coding tasks among others. The use of benchmarks allows developers to fine-tune a model, and a user to choose a model that better fits their needs or adjust its responses by changing a system prompt and parameters like temperature. This requires a lot of benchmark runs. Running a benchmark is costly, time- and energy-consuming. Also, it is often inefficient, as they may contain problems that are not relevant for the use case of a model being adjusted.

So the goal of this work is to explore existing benchmarks, and design and develop an approach to benchmarking that is highly customizable and extendible, time-efficient, eco-friendly. As a result, a modular benchmark was developed, that allows to add and modify test cases using a user interface, and configure a benchmark run using task filters and toggles for enabling or disabling specific checks. Also, the output of the benchmark allows to examine the results in detail.

**Keywords:** keyword1, keyword2, ..., keyword5





---

# Índice de contenidos

---

<b>1. Introduction</b>	<b>1</b>
1.1. Problem Statement . . . . .	2
1.2. Objectives . . . . .	2
<b>2. State of the Art</b>	<b>3</b>
2.1. Evolution of Code Generation Benchmarks . . . . .	3
2.2. Types of LLM Benchmarks . . . . .	4
2.3. Problems in LLM Benchmarking . . . . .	5
2.4. Environmental impact of Benchmarking . . . . .	5
2.5. Existing Benchmarking Frameworks and Their Limitations . . . . .	6
<b>3. Problem Description</b>	<b>7</b>
3.1. Analysis of Current Limitations . . . . .	7
3.2. Requirements for a Solution . . . . .	7
<b>4. Proposed Solution</b>	<b>9</b>
4.1. Architecture Overview . . . . .	9
4.2. Task Dataset Format . . . . .	9
4.3. Interactive Configuration . . . . .	9
<b>5. Implementation Details</b>	<b>11</b>
5.1. Project Structure . . . . .	11
5.2. Backend Development . . . . .	11
5.2.1. REST API Endpoints . . . . .	11
5.2.2. Core Functionality . . . . .	12
5.3. Frontend Development . . . . .	12
5.3.1. User Interface Components . . . . .	12
5.3.2. Interactive Features . . . . .	12
5.4. Deployment . . . . .	12
5.4.1. Local Development . . . . .	13
5.4.2. Docker Deployment . . . . .	13
<b>6. Results and Evaluation</b>	<b>15</b>
6.1. Performance Analysis . . . . .	15
6.2. Case Studies . . . . .	15
<b>7. Conclusions</b>	<b>17</b>
7.1. Contributions . . . . .	17

## ÍNDICE DE CONTENIDOS

---

7.2. Future Work . . . . .	17
7.3. Impact . . . . .	17
<b>Bibliografía</b>	<b>19</b>
<b>Apéndice A. Ejemplo de anexo</b>	<b>23</b>





# CAPÍTULO 1

---

## Introduction

---

Large Language Models (LLMs) have revolutionized software development through AI-assisted programming tools like GitHub Copilot. However, evaluating and fine-tuning these models requires extensive benchmarking, which comes at significant computational, financial, and environmental costs. Current benchmarks often contain irrelevant tasks and provide limited customization options, making them inefficient for specific use cases.

This work explores existing benchmarks for code-generating LLMs and proposes a novel approach: modular and customizable benchmarks that can be tailored to specific needs while remaining cost-effective and environmentally conscious.

This thesis began with a comprehensive literature review, using recent critical reviews as a foundation for analyzing the state of LLM benchmarking. The review was extended by following citations to relevant articles published in late 2024 and 2025, focusing on keywords related to benchmarking and LLMs. The findings from this literature analysis informed the design and implementation of a modular benchmarking system, which was then evaluated for efficiency, flexibility, and environmental impact.

The literature review and subsequent research were guided by the following questions:

- What are the main limitations of current LLM benchmarks for code generation?
- What metrics best reflect real-world usability and code quality?
- How can benchmarks be made customizable for different user needs?
- What is the environmental impact of repeated benchmarking, and how can it be reduced?

### **Main contributions of this work:**

- A critical analysis of existing LLM code generation benchmarks and their shortcomings.
- The design and implementation of a modular, customizable benchmarking framework.
- Integration of environmental and cost considerations into the benchmarking process.
- An interactive web interface for configuring benchmarks and analyzing results.

- Recommendations for future benchmarking practices based on empirical findings.

### 1.1. Problem Statement

Running benchmarks is essential for LLM developers, researchers and users. Developers and researchers need benchmarks to fine-tune their models or compare approaches, while users rely on them to select appropriate models and adjust parameters like temperature and system prompts. However, this process faces several challenges: **Cost and Resource Consumption:** Each benchmark run consumes significant computational resources, time, and energy. As noted in recent studies, the environmental impact of repeated benchmark runs is becoming a growing concern.

**Lack of Customization:** Most benchmarks are developed as fixed sets of tasks, making them difficult to adapt for specific use cases. This limitation was highlighted in the SWE-bench study [Jimenez \*et al.\* 2024](#), which emphasized the need for more flexible evaluation methods.

**Inefficient Task Selection:** Many benchmarks include tasks that may be irrelevant for specific applications, leading to wasted resources. Recent research [Vendrow \*et al.\* 2025](#) has shown that benchmark saturation often results in running unnecessary tests that all modern LLMs pass easily.

### 1.2. Objectives

The main objectives of this work are:

1. Analyze existing benchmarks and their limitations
2. Design a modular benchmark system that allows:
  - Custom task selection and filtering
  - Configuration of testing criteria
  - Support for multiple programming languages and task types
  - Integration with CI/CD pipelines
3. Implement an interactive web interface for benchmark configuration and result analysis
4. Develop a cost-efficient and environmentally conscious approach to benchmarking

# CAPÍTULO 2

---

## State of the Art

---

### 2.1. Evolution of Code Generation Benchmarks

The evolution of benchmarks for code generation has been driven by the need to evaluate the capabilities of Large Language Models (LLMs) in programming tasks. Early benchmarks focused on isolated tasks, but as LLMs became more sophisticated, the need for more comprehensive and realistic evaluation methods emerged.

The first benchmarks were aimed at text comprehension and contained questions and expected answers, such as GLUE, SQuAD and GSM8K with grade-school math problems ([Vendrow et al. 2025](#)).

As LLM capabilities expanded, benchmarks shifted towards programming tasks, with a focus on code generation and understanding. Some pioneers in LLM code benchmarking that are MBPP, HumanEval and APPS.

- MBPP (Mostly Basic Python Problems) published by [Austin et al. 2021](#) contains 974 crowdsourced Python programming problems with tests.
- HumanEval was developed in OpenAI by [Chen et al. 2021](#) with 164 hand-crafted problems, with a goal of avoiding data leakage (to ensure that the problems and golden solutions are not present in the training dataset).
- APPS by [Hendrycks et al. 2021](#) features 10000 Python tasks with 131777 test cases, borrowed from open-access sites like Codewars and Codeforces.

These benchmarks are still used to this day for comparing performance of different models in scientific papers. Notably, APPS is a benchmark commonly used for fine-tuning LLMs for programming tasks, as it allows to use separate sets of problems for training and evaluation ([Ben Allal et al. 2022](#)).

The mentioned benchmarks became less effective, as newer models were trained on the same tasks they are being evaluated on. This phenomenon is known as **data leakage** as described by [Vendrow et al. 2025](#).

Amazon's Recode benchmark [Wang et al. 2022](#) addressed this issue by introducing perturbations on docstrings, function names, and codes, while staying semantically close



to the original task. However, that is more of a way to test the robustness of the model, rather than a way to test its ability to solve brand-new problems.

More recent developments like HumanEval Pro and MBPP Pro [Yu et al. 2024](#) introduced more sophisticated testing approaches. Their multistep evaluation process tests an LLM's ability to work with its own generated code. First, an LLM generates a solution to a known problem from HumanEval or MBPP datasets. Then, it is given a new task that requires calling a function generated in the first step. This approach revealed limitations in some models that perform well on simpler tasks.

While academic benchmarks like HumanEval focus on controlled, isolated tasks, initiatives like SWE-bench [Jimenez et al. 2024](#) have moved toward real-world scenarios. Software engineering tasks were taken from resolved Issues from GitHub repositories of open-source projects Python. To test the generated solution, the authors identified the test cases that were failing before the solution was applied, and then checked if the solution passed all the tests. SWE-bench is famous for its leaderboards, where laboratories and companies all over the world compete to achieve the higher percentage of solved tasks.

Researchers [Chi et al. 2025](#) have found a different way to evaluate LLMs in real-world scenarios. Instead of using a fixed set of tasks, they developed a plugin called CopilotArena for an IDE. The plugin provides two code completion options from different LLMs, and allows a user to choose the one they prefer. This approach allows for a more realistic evaluation of LLMs in coding tasks, but it lacks the controlled environment and a solid numerical result for each model. Such an approach could be useful for A/B testing of LLMs in production, but it is not suitable for scientific research and repeated evaluations during fine-tuning.

The benchmarks mentioned above perform in a static environment, where the model is given a task and expected to generate a solution. The InterCode framework by [Yang et al. 2023](#) introduces interactive environments using Docker. This enables evaluation of LLMs in realistic and interactive development scenarios with compilation and runtime feedback. The environments and scenarios were prepared for Python, SQL, and BASH, but the framework allows introducing new environments and scenarios. This approach more closely mirrors actual developer workflows and allows for testing LLMs in the role of a partially independent agent. But this approach comes with increased computational overhead of running a Docker environment, a virtual operating system, and an instance of a database, which limits the overall speed and a number of scenarios that can be tested at once.

## 2.2. Types of LLM Benchmarks

In [Chi et al. 2025](#) the authors developed a plugin for IDE that offers the user two code completion options from two different LLMs and records which option the user chose. This allowed them to compare the performance of different LLMs in real-world coding tasks, providing valuable insights into user preferences and model effectiveness. Another article [Mozannar et al. 2024](#) finds out that a user can often accept a first proposed solution in order to see it with proper syntax highlighting and be able to understand it better. But later he could even remove the suggestion and waits for a repeated completion or

simply writes the piece of the code themselves. Although that should not be considered a significant flaw of the benchmark, it could partially skew the results.

Most benchmarks are very limited in the output they provide. For example, SWE leaderboard [SWE-bench Leaderboards s.f.](#) is created based on a single number that does not reflect the types of tasks that the model is better or worse at solving [Miah y Zhu 2024](#). Thus, a user might choose a suboptimal model for their specific needs, resulting in lower performance or higher cost. This is partially countered by websites that aggregate results on several benchmarks [LLM Leaderboard 2025 — vellum.ai s.f.](#) which can give a very high-level picture.

## 2.3. Problems in LLM Benchmarking

Benchmark saturation mentioned above is partially explained by advances in models, but it can also be attributed to information **leaking**: the popular and publicly available benchmarks appear in the training datasets accompanied with the golden solutions. Even then, it doesn't mean that an LLM won't struggle when presented with the same task. When changing the task phrasing while keeping the semantic consistent, there is a 4.5 percent drop in solvability, showing that the models remember the phrasing of the descriptions in the original dataset. [Uniyal et al. 2024](#)

One of the important aspects of LLM evaluation is the choice of the metrics. For code quality, there are BLEU, CodeBLEU, RUBY, ROUGE-L, METEOR, ChrF. They assess the similarity of the generated code to the golden solution, taking into account the properties of source code. Evtikhiev et al. [Evtikhiev et al. 2023](#) takes 6 metrics, commonly used in papers. However, the authors conduct a study, comparing the results of metrics with human evaluation of the solutions. The results suggest that none of the analyzed metrics can correctly emulate human judgment, but ChrF metric is considered better than the others commonly used in papers.

A paper by [Crupi et al. 2025](#) looks into an approach of using LLM to evaluate the quality of the solution generated by another model (LLM-as-a-judge approach). As a result, they come to a conclusion that LLM-as-a-judge is a substantial improvement over mentioned metrics, and GPT-4-turbo can mimic closely a human evaluation.

## 2.4. Environmental impact of Benchmarking

Repeated training and benchmarking of LLMs require substantial computational resources, leading to significant electricity consumption and carbon emissions. This environmental impact is increasingly important in the context of global efforts to reduce carbon footprints. We will want for benchmarks to account for these factors and encourage more sustainable evaluation practices.

There are leaderboards that account for  $CO_2$  emissions, such as Hugging Face [LLM CO2 emissions calculation - a Hugging Face Space — huggingface.co s.f.](#), which tracks the carbon footprint of models. However, these metrics are often not integrated into

traditional benchmarks, leading to a lack of awareness about the environmental impact of LLM evaluation practices.

The most common metric in benchmarks is `pass@k` that measures the percentage of correct solutions among the  $k$  solutions generated by the model. This implies that for each task in the benchmark dataset, a model repeatedly generates a number of solutions, just to receive a single numeric result to use for a metric. This metric is used in ClassEval, MBPP, MathQA-Python, CoderEval, and HumanEval+. Notably, HumanEval and HumanEval+ use  $k = 100$  (`pass@100`). However, as Miah and Zhu [Miah y Zhu 2024](#) pointed out, users do not normally run the LLM several times, so `pass@k` does not reflect its usability.

## 2.5. Existing Benchmarking Frameworks and Their Limitations

Apart from the benchmarks themselves, there are several frameworks that facilitate LLM evaluation. These frameworks provide tools for running benchmarks and collecting results.

Two widely used benchmarking frameworks are **bigcode-evaluation-harness** [Ben Allal et al. 2022](#) and **lm-evaluation-harness** [Gao et al. 2024](#). Both provide tools for running standardized benchmarks on LLMs, but they have notable limitations. The lm-evaluation-harness is more general-purpose and supports a broader range of language tasks, yet it also relies on fixed task sets and lacks modularity for user-defined benchmarks. Neither framework provides built-in support for environmental metrics or fine-grained task selection, highlighting the need for more flexible and sustainable benchmarking solutions.

Feature	bigcode-evaluation-harness	lm-evaluation-harness
Evaluation of multiple LLMs	One model per run	One model per run
Available configuration	Task dataset name, Number of tasks, Temperature, Saving LLM responses, Limit of LLM response	TO-DO
Included benchmarks	MBPP(+), DS-1000, MultiPL-E, Mercury, etc.	Limited, fixed task sets
Defining new tasks	Requires source code modification	Requires source code modification
Run interface	CLI-based, no GUI	CLI-based, no GUI
Result analysis	Overall numeric metric and LLM responses saved as a file	
Visualization	No visualization tools	No visualization tools

Tabla 2.1: Comparison of bigcode-evaluation-harness and lm-evaluation-harness

# CAPÍTULO 3

---

## Problem Description

---

### 3.1. Analysis of Current Limitations

The limitations of existing benchmarks extend beyond just inefficiency. Studies like have identified several critical issues:

- Benchmark saturation: Many tasks become irrelevant as models improve [Vendrow et al. 2025](#)
- Task leaking: LLMs are trained on benchmark tasks solutions, making them obsolete [Vendrow et al. 2025](#)
- Limited feedback: Most benchmarks provide only pass/fail results
- High costs: Running comprehensive benchmarks is expensive and time-consuming
- Environmental impact: Repeated runs contribute to unnecessary carbon emissions
- Inflexibility: Fixed task sets don't adapt to specific needs
- Error-prone: benchmarks contain up to 5 percent of mislabelled or erroneous tasks [Vendrow et al. 2025](#)

### 3.2. Requirements for a Solution

Based on the analyzed limitations and user needs, we identified key requirements for a more effective benchmarking approach:

Modularity:

- Support for multiple programming languages
- Ability to add/modify tasks easily
- Configurable testing criteria

Efficiency:

- Task filtering capabilities

- Optimized resource usage
- Quick feedback loops

User Experience:

- Interactive configuration interface
- Detailed result analysis
- Integration with development workflows

The benchmark should:

- Provide a way to analyse individual task failures
- Allow users to select tasks based on their specific needs
- Get the most information from each generated solution
- ...

# CAPÍTULO 4

---

## Proposed Solution

---

### 4.1. Architecture Overview

Our solution combines three main components:

- A Spring Boot backend with MVC architecture
- A React-based frontend for configuration and visualization
- Docker environments for isolated task execution

Figura 4.1: System Architecture Overview

### 4.2. Task Dataset Format

Tasks are defined in YAML format, allowing for easy modification and extension:

```
tasks:
- name: "Example Task"
  type: "implementation"
  difficulty: "medium"
  languages: ["python", "java"]
  parameters:
    use_libraries: false
    generate_tests: true
```

### 4.3. Interactive Configuration

The React frontend provides: - Task filtering by type, difficulty, and language - Testing criteria selection - Resource usage configuration - Result visualization



# CAPÍTULO 5

---

## Implementation Details

---

### 5.1. Project Structure

The project follows a modern microservices architecture with separate backend and frontend applications. Currently, the configuration, status and result files are stored in a file-based storage system, but they can be easily adapted to use a database if needed.

Project Root

```
|-- src/           # Spring Boot application source code
|-- frontend/      # React application source code
|-- exec_configs/  # Execution configuration YAML files
|-- task_sources/  # Task source YAML files
|-- bench_status/  # Benchmark status files
'-- bench_results/ # Benchmark results
```

### 5.2. Backend Development

The Spring Boot application serves as the core of the system, providing:

#### 5.2.1. REST API Endpoints

The backend exposes a comprehensive RESTful API:

- `/api/configs`: Configuration file management
- `/api/tasks`: Task source file operations
- `/api/benchmarks`: Benchmark execution control
- `/api/status`: Status monitoring for current benchmarks run
- `/api/results`: Result retrieval for visualization and analysis



### 5.2.2. Core Functionality

- Task processing and execution in isolated environments
- Docker container management for language-specific runtimes
- File-based storage system for configurations and results
- Integration with LLM judges through Spring AI

Figura 5.1: Backend Component Architecture

## 5.3. Frontend Development

The React frontend, built with TypeScript, provides an intuitive interface for:

### 5.3.1. User Interface Components

- Configuration file upload and management
- Task source file handling
- Real-time benchmark monitoring
- Result visualization and analysis

### 5.3.2. Interactive Features

The application enables users to:

- Upload and manage YAML configurations
- Configure benchmark parameters interactively
- Monitor benchmark progress in real-time
- Analyze and export results

Figura 5.2: Frontend Interface Overview

## 5.4. Deployment

Both components support flexible deployment options:

### 5.4.1. Local Development

```
# Backend
mvn spring-boot:run
```

```
# Frontend
cd frontend
npm install
npm run dev
```

### 5.4.2. Docker Deployment

The system includes Docker configurations for containerized deployment:

```
# Backend container
docker build -t ai-benchmark-backend -f Dockerfile.backend .
docker run -p 8080:8080 ai-benchmark-backend

# Frontend container
docker build -t ai-benchmark-frontend -f Dockerfile.frontend .
docker run -p 5173:5173 ai-benchmark-frontend
```



# CAPÍTULO 6

---

## Results and Evaluation

---

### 6.1. Performance Analysis

We evaluated our solution against traditional benchmarks:

Tabla 6.1: Comparison with Traditional Benchmarks

Key metrics:

- Execution time per task
- Resource usage (CPU, memory)
- Cost per benchmark run
- Environmental impact (estimated  $CO_2$  emissions)

### 6.2. Case Studies

We present three scenarios demonstrating the system's flexibility:

1. Fine-tuning an LLM for specific programming tasks
2. Evaluating code style consistency
3. Testing real-world project integration



# CAPÍTULO 7

---

## Conclusions

---

### 7.1. Contributions

Our modular benchmark approach offers several advantages:

- Customizable evaluation scenarios
- Reduced resource consumption
- Improved feedback quality
- Easy integration with development workflows

### 7.2. Future Work

Potential improvements include:

- Additional programming language support
- Enhanced visualization tools
- CI/CD pipeline integration
- Advanced metrics collection

### 7.3. Impact

This work contributes to more sustainable and efficient LLM evaluation practices, potentially reducing both costs and environmental impact while providing more meaningful results.



---

## Bibliografía

---

1. C. E. Jimenez *et al.*, *SWE-bench: Can Language Models Resolve Real-World GitHub Issues?*, 2024, arXiv: 2310.06770 [cs.CL], (<https://arxiv.org/abs/2310.06770>).
2. J. Vendrow, E. Vendrow, S. Beery y A. Madry, *Do Large Language Model Benchmarks Test Reliability?*, 2025, arXiv: 2502.03461 [cs.LG], (<https://arxiv.org/abs/2502.03461>).
3. J. Austin *et al.*, Program Synthesis with Large Language Models. *arXiv preprint arXiv:2108.07732* (2021).
4. M. Chen *et al.*, *Evaluating Large Language Models Trained on Code*, 2021, arXiv: 2107.03374 [cs.LG], (<https://arxiv.org/abs/2107.03374>).
5. D. Hendrycks *et al.*, Measuring Coding Challenge Competence With APPS. *NeurIPS* (2021).
6. L. Ben Allal *et al.*, *A framework for the evaluation of code generation models*, <https://github.com/bigcode-project/bigcode-evaluation-harness>, 2022.
7. S. Wang *et al.*, ReCode: Robustness Evaluation of Code Generation Models. (<https://arxiv.org/abs/2212.10264>) (2022).
8. Z. Yu, Y. Zhao, A. Cohan y X.-P. Zhang, *HumanEval Pro and MBPP Pro: Evaluating Large Language Models on Self-invoking Code Generation*, 2024, arXiv: 2412.21199 [cs.SE], (<https://arxiv.org/abs/2412.21199>).
9. W. Chi *et al.*, *Copilot Arena: A Platform for Code LLM Evaluation in the Wild*, 2025, arXiv: 2502.09328 [cs.SE], (<https://arxiv.org/abs/2502.09328>).
10. J. Yang, A. Prabhakar, K. Narasimhan y S. Yao, *InterCode: Standardizing and Benchmarking Interactive Coding with Execution Feedback*, 2023, arXiv: 2306.14898 [cs.CL], (<https://arxiv.org/abs/2306.14898>).
11. H. Mozannar, G. Bansal, A. Fourney y E. Horvitz, *Reading Between the Lines: Modeling User Behavior and Costs in AI-Assisted Programming*, 2024, arXiv: 2210.14306 [cs.SE], (<https://arxiv.org/abs/2210.14306>).
12. *SWE-bench Leaderboards*, <https://www.swebench.com/>, [Accessed 17-08-2025].
13. T. Miah y H. Zhu, *User Centric Evaluation of Code Generation Tools*, 2024, arXiv: 2402.03130 [cs.SE], (<https://arxiv.org/abs/2402.03130>).
14. *LLM Leaderboard 2025 — vellum.ai*, <https://www.vellum.ai/llm-leaderboard>, [Accessed 17-08-2025].
15. M. Uniyal *et al.*, presentado en Findings of the Association for Computational Linguistics: EMNLP 2024, págs. 15397-15402.



16. M. Evtikhiev, E. Bogomolov, Y. Sokolov y T. Bryksin, Out of the bleu: how should we assess quality of the code generation models? *Journal of Systems and Software* **203**, 111741 (2023).
17. G. Crupi *et al.*, On the Effectiveness of LLM-as-a-judge for Code Generation and Summarization. *IEEE Transactions on Software Engineering* (2025).
18. *LLM CO2 emissions calculation - a Hugging Face Space* — *huggingface.co*, [https://huggingface.co/docs/leaderboards/open\\_llm\\_leaderboard/emissions](https://huggingface.co/docs/leaderboards/open_llm_leaderboard/emissions), [Accessed 17-08-2025].
19. L. Gao *et al.*, *The Language Model Evaluation Harness*, ver. v0.4.3, jul. de 2024, (<https://zenodo.org/records/12608602>).

## Apéndice



## APÉNDICE A

---

### Ejemplo de anexo

---

Anexo primero.